



1 Three chat conditions

HH real



HH roleplay



H-LLM roleplay



GPT-OSS Llama 3.3 Mixtral

three model families, 27 case studies

535 conversations, 29,971 spans

2 Learned segmentation

learned

1

2

3

regex

adapter never sees study data

ordering robust under three units

3 Interaction coding

OnCoCo classifier



38 counselor 28 client

one classifier for all conditions

4 Anchored comparison

the observed distance



HH real H-LLM

what chance alone produces

split-half band



→ JSD

HH real's 53 client-side conversations cut in two disjoint halves


size-matched null



→ JSD

both groups redrawn from all 53, with repeats, at the sizes really compared

size-matched P95 0.073



split-half P95 0.132

client side throughout; Cramér's V reported alongside as effect size

HH roleplay

0.233

H-LLM

0.453

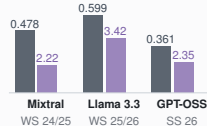
4.2× beyond the noise floor

JSD is 0 when two label mixes are identical and grows as they part. It brings no scale of its own.

5 Convergent validity

measured distance

perceived



one course line, three semesters, client deployment exchanged

perception tracks the measurement